


# *Machine Learning Model*

To Car and Heart Datasets



Esakiappan E

20 September, 2024

---

# Machine Learning Algorithm :

## Regressor :

- Linear Regressor
- Ridge Regressor
- Lasso Regressor




## Classifier :

- Logistic Classification
- Naive Bayes Classifier

## For Both,

- Decision Tree Algorithm
- KNN Algorithm
- Random Forest Algorithm



# Regressor(Car Dataset)

Total Rows and columns  
2059 x 20



	Make	Model	Price	Year	Kilometer	Fuel Type	Transmission	Location	Color	Owner	Seller Type	Engine	Max Power	Max Torque	Drivetrain	Length
0	Honda	Amaze 1.2 VX i-VTEC	505000	2017	87150	Petrol	Manual	Pune	Grey	First	Corporate	1198 cc	87 bhp @ 6000 rpm	109 Nm @ 4500 rpm	FWD	3990.0
1	Maruti Suzuki	Swift DZire VDI	450000	2014	75000	Diesel	Manual	Ludhiana	White	Second	Individual	1248 cc	74 bhp @ 4000 rpm	190 Nm @ 2000 rpm	FWD	3995.0
2	Hyundai	i10 Magna 1.2 Kappa2	220000	2011	67000	Petrol	Manual	Lucknow	Maroon	First	Individual	1197 cc	79 bhp @ 6000 rpm	112.7619 Nm @ 4000 rpm	FWD	3585.0
3	Toyota	Glanza G	799000	2019	37500	Petrol	Manual	Mangalore	Red	First	Individual	1197 cc	82 bhp @ 6000 rpm	113 Nm @ 4200 rpm	FWD	3995.0
4	Toyota	Innova 2.4 VX 7 STR [2016-2020]	1950000	2018	69000	Diesel	Manual	Mumbai	Grey	First	Individual	2393 cc	148 bhp @ 3400 rpm	343 Nm @ 1400 rpm	RWD	4735.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
2054	Mahindra	XUV500 W8 [2015-2017]	850000	2016	90300	Diesel	Manual	Surat	White	First	Individual	2179 cc	138 bhp @ 3750 rpm	330 Nm @ 1600 rpm	FWD	4585.0
2055	Hyundai	Eon D-Lite +	275000	2014	83000	Petrol	Manual	Ahmedabad	White	Second	Individual	814 cc	55 bhp @ 5500 rpm	75 Nm @ 4000 rpm	FWD	3495.0
2056	Ford	Figo Duratec Petrol ZXI 1.2	240000	2013	73000	Petrol	Manual	Thane	Silver	First	Individual	1196 cc	70 bhp @ 6250 rpm	102 Nm @ 4000 rpm	FWD	3795.0
2057	BMW	5-Series 520d Luxury Line [2017-2019]	4290000	2018	60474	Diesel	Automatic	Coimbatore	White	First	Individual	1995 cc	188 bhp @ 4000 rpm	400 Nm @ 1750 rpm	RWD	4936.0
2058	Mahindra	Bolero Power Plus ZLX [2016-2019]	670000	2017	72000	Diesel	Manual	Guwahati	White	First	Individual	1493 cc	70 bhp @ 3600 rpm	195 Nm @ 1400 rpm	RWD	3995.0

2059 rows × 20 columns



# Libraries

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2059 entries, 0 to 2058
Data columns (total 20 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   Make                2059 non-null  object 
 1   Model               2059 non-null  object 
 2   Price               2059 non-null  int64  
 3   Year               2059 non-null  int64  
 4   Kilometer           2059 non-null  int64  
 5   Fuel Type           2059 non-null  object 
 6   Transmission        2059 non-null  object 
 7   Location            2059 non-null  object 
 8   Color              2059 non-null  object 
 9   Owner               2059 non-null  object 
10   Seller Type         2059 non-null  object 
11   Engine              1979 non-null  object 
12   Max Power           1979 non-null  object 
13   Max Torque          1979 non-null  object 
14   Drivetrain          1923 non-null  object 
15   Length              1995 non-null  float64 
16   Width               1995 non-null  float64 
17   Height              1995 non-null  float64 
18   Seating Capacity    1995 non-null  float64 
19   Fuel Tank Capacity  1946 non-null  float64 
dtypes: float64(5), int64(3), object(12)
memory usage: 321.8+ KB
```

```
import sklearn
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
```

# Data Preprocessing

## Missing Value Imputation

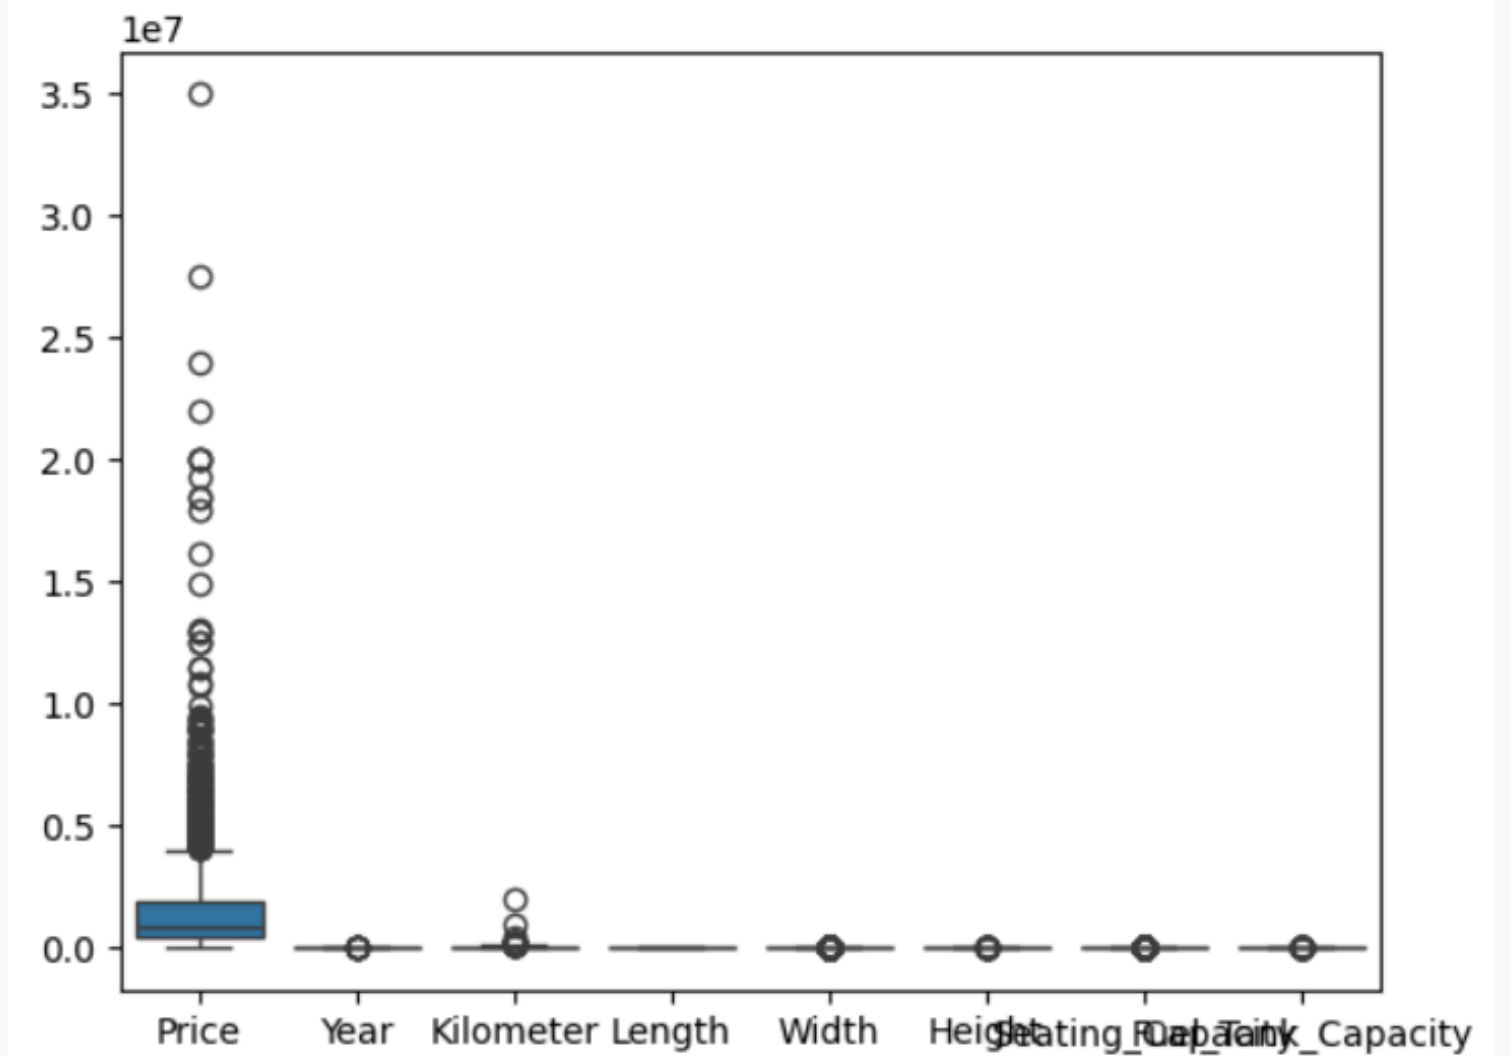
```
In [224]: df.isnull().sum()
```

```
Out[224]: Make          0
Price          0
Year           0
Kilometer      0
Fuel Type      0
Transmission   0
Location       0
Color          0
Owner          0
Seller Type    0
Engine         80
Max Power      80
Max Torque     80
Drivetrain     136
Length         64
Width          64
Height         64
Seating Capacity 64
Fuel Tank Capacity 113
dtype: int64
```

## Outliers

```
sns.boxplot(df)
```

<Axes: >





# After Data Preprocessing

Total Rows and columns  
1290 x 18

14]: df

14]:


	Make	Price	Year	Kilometer	Fuel_Type	Transmission	Color	Owner	Seller_Type	Engine	Max_Power	Max_Torque	Drivetrain	Length	Width	H
0	Honda	505000	2017	87150	Petrol	Manual	Grey	First	Corporate	1198 cc	87 bhp @ 6000 rpm	109 Nm @ 4500 rpm	FWD	3990.0	1680.0	1
1	Maruti Suzuki	450000	2014	75000	Diesel	Manual	White	Second	Individual	1248 cc	74 bhp @ 4000 rpm	190 Nm @ 2000 rpm	FWD	3995.0	1695.0	1
2	Hyundai	220000	2011	67000	Petrol	Manual	Maroon	First	Individual	1197 cc	79 bhp @ 6000 rpm	112.7619 Nm @ 4000 rpm	FWD	3585.0	1595.0	1
3	Toyota	799000	2019	37500	Petrol	Manual	Red	First	Individual	1197 cc	82 bhp @ 6000 rpm	113 Nm @ 4200 rpm	FWD	3995.0	1745.0	1
5	Maruti Suzuki	675000	2017	73315	Petrol	Manual	Grey	First	Individual	1373 cc	91 bhp @ 6000 rpm	130 Nm @ 4000 rpm	FWD	4490.0	1730.0	1
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
2051	Maruti Suzuki	925000	2021	48000	Petrol	Manual	White	First	Individual	1462 cc	103 bhp @ 6000 rpm	138 Nm @ 4400 rpm	FWD	3995.0	1790.0	1
2052	Hyundai	409999	2014	68000	Diesel	Manual	Silver	First	Individual	1396 cc	90@4000	220@1750	FWD	3940.0	1710.0	1
2053	Maruti Suzuki	245000	2014	79000	Petrol	Manual	White	Second	Individual	1197 cc	85 bhp @ 6000 rpm	113 Nm @ 4500 rpm	FWD	3775.0	1680.0	1
2055	Hyundai	275000	2014	83000	Petrol	Manual	White	Second	Individual	814 cc	55 bhp @ 5500 rpm	75 Nm @ 4000 rpm	FWD	3495.0	1550.0	1
2056	Ford	240000	2013	73000	Petrol	Manual	Silver	First	Individual	1196 cc	70 bhp @ 6250 rpm	102 Nm @ 4000 rpm	FWD	3795.0	1680.0	1

1290 rows × 18 columns



# Heatmap

xes: >



# Anova Testing

```
from sklearn.feature_selection import f_classif
a=f_classif(x,y)
a
```

```
C:\Users\DELL\anaconda3\Lib\site-packages\sklearn\feature_selection\_univariate_constant:
warnings.warn("Features %s are constant." % constant_features_idx, UserWarning)
```

```
C:\Users\DELL\anaconda3\Lib\site-packages\sklearn\feature_selection\_univariate_constant:
encountered in divide
f = msb / msw
```

```
(array([1.24270134, 3.31685859, 1.26020872, 1.26666946, 2.96390008,
        1.15363953, 1.36857749, 1.45463052, 1.73468918, 1.65943316,
        1.83132598, 2.22699034, 2.92017757, 4.80012178, 1.73410167,
        nan, 3.19883692]),
 array([4.54330865e-03, 4.34892377e-50, 2.73358074e-03, 2.25497778e-03,
        2.31083587e-41, 4.34007910e-02, 7.79876904e-05, 2.98218367e-06,
        1.02852885e-11, 3.83068584e-10, 8.13918129e-14, 4.45838462e-23,
        2.83451528e-40, 1.51808401e-84, 1.05857104e-11, nan,
        3.47367394e-47]))
```

```
: a=pd.Series(a[1])
a.index=x.columns
a
```

```
: Make      4.543309e-03
Year      4.348924e-50
Kilometer  2.733581e-03
Fuel_Type  2.254978e-03
Transmission 2.310836e-41
Color      4.340079e-02
Owner      7.798769e-05
Seller_Type 2.982184e-06
Engine     1.028529e-11
Max_Power  3.830686e-10
Max_Torque 8.139181e-14
Drivetrain 4.458385e-23
Length     2.834515e-40
Width      1.518084e-84
Height     1.058571e-11
Seating_Capacity NaN
Fuel_Tank_Capacity 3.473674e-47
dtype: float64
```





# Support vector Machine

## Random Forest Regressor

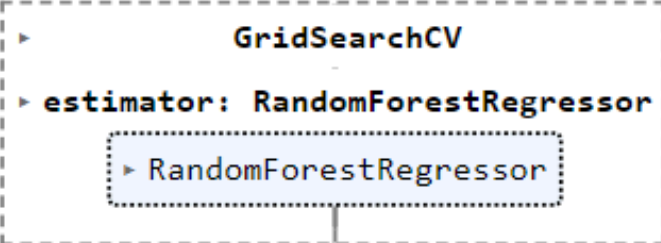
```
In [117]: from sklearn.ensemble import RandomForestRegressor
```

```
In [118]: rf = RandomForestRegressor()
```

```
In [119]: param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [3, 5, 7, 9],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}
```

```
In [120]: grid_search = GridSearchCV(RandomForestRegressor(), param_grid, cv=5, scoring='neg_mean_squared_error')
grid_search.fit(x_train, y_train)
```

```
Out[120]:
```



```
  > GridSearchCV
  > estimator: RandomForestRegressor
    > RandomForestRegressor
```

```
In [121]: grid_search.best_params_
```

```
Out[121]: {'max_depth': 9,
    'min_samples_leaf': 1,
    'min_samples_split': 5,
    'n_estimators': 200}
```

```
In [122]: grid_search.best_score_
```

```
Out[122]: -82303062247.39612
```



```
In [123]: best_rf = grid_search.best_estimator_  
best_rf.fit(x_train, y_train)
```

```
Out[123]: 

RandomForestRegressor  
RandomForestRegressor(max_depth=9, min_samples_split=5, n_estimators=200)


```

```
In [124]: y_pre = best_rf.predict(x_test)
```

```
In [125]: mse = mean_squared_error(y_test, y_pre)  
mse
```

```
Out[125]: 82589825641.39249
```

```
In [126]: mae = mean_absolute_error(y_test, y_pre)  
mae
```

```
Out[126]: 150024.63709946853
```


```
In [127]: r2 = r2_score(y_test, y_pre)  
r2
```

```
Out[127]: 0.8798203930254364
```



# Classifier(Heart Dataset)

Total Rows and columns  
1025 x 14



	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	0
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	0
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	0
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
1020	59	1	1	140	221	0	1	164	1	0.0	2	0	2	1
1021	60	1	0	125	258	0	0	141	1	2.8	1	1	3	0
1022	47	1	0	110	275	0	0	118	1	1.0	1	1	2	0
1023	50	0	0	110	254	0	0	159	0	0.0	2	0	2	1
1024	54	1	0	120	188	0	1	113	0	1.4	1	1	3	0

1025 rows × 14 columns



## Libraries

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2059 entries, 0 to 2058
Data columns (total 20 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   Make                  2059 non-null   object  
 1   Model                 2059 non-null   object  
 2   Price                 2059 non-null   int64   
 3   Year                  2059 non-null   int64   
 4   Kilometer             2059 non-null   int64   
 5   Fuel Type             2059 non-null   object  
 6   Transmission          2059 non-null   object  
 7   Location              2059 non-null   object  
 8   Color                 2059 non-null   object  
 9   Owner                 2059 non-null   object  
10   Seller Type           2059 non-null   object  
11   Engine                1979 non-null   object  
12   Max Power             1979 non-null   object  
13   Max Torque            1979 non-null   object  
14   Drivetrain            1923 non-null   object  
15   Length                1995 non-null   float64  
16   Width                 1995 non-null   float64  
17   Height                1995 non-null   float64  
18   Seating Capacity      1995 non-null   float64  
19   Fuel Tank Capacity    1946 non-null   float64  
dtypes: float64(5), int64(3), object(12)
memory usage: 321.8+ KB
```

```
import sklearn
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
```



# Data Preprocessing

## Missing Value Imputation


```
] df.isnull().sum()
```

```
] age      0  
sex       0  
cp        0  
trestbps  0  
chol      0  
fbs       0  
restecg   0  
thalach   0  
exang     0  
oldpeak   0  
slope     0  
ca        0  
thal      0  
target    0  
dtype: int64
```

## Outliers

```
sns.boxplot(df)
```

<Axes: >



# After Data Preprocessing

Total Rows and columns  
**964 x 14**


df														
	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	0
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	0
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	0
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
1020	59	1	1	140	221	0	1	164	1	0.0	2	0	2	1
1021	60	1	0	125	258	0	0	141	1	2.8	1	1	3	0
1022	47	1	0	110	275	0	0	118	1	1.0	1	1	2	0
1023	50	0	0	110	254	0	0	159	0	0.0	2	0	2	1
1024	54	1	0	120	188	0	1	113	0	1.4	1	1	3	0

964 rows × 14 columns


# Univariate Analysis

## Histogram


```
In [16]: plt.hist(df['age'],bins=10,color='green')
Out[16]: (array([ 4., 39., 109., 125., 117., 187., 203., 128., 43., 9.]),
array([29. , 33.8, 38.6, 43.4, 48.2, 53. , 57.8, 62.6, 67.4, 72.2, 77. ]))
<BarContainer object of 10 artists>
```



## Bar Chart



# Heatmap





# Anova Testing

```
20]: from sklearn.feature_selection import f_classif
a=f_classif(x,y)
a
```

```
20]: (array([ 52.55650761, 108.13170468, 205.80610771,  11.77302565,
            15.78805862,   1.52622117,  29.0802467 , 213.91321996,
            228.44290963, 239.65858734, 116.19439063, 156.88734103,
            132.67843121]),
      array([8.58053462e-13, 4.48707752e-24, 1.93641102e-42, 6.26516229e-04,
            7.61409501e-05, 2.16982079e-01, 8.74096898e-08, 6.83925686e-44,
            1.81223443e-46, 1.95455095e-48, 1.17560747e-25, 1.88403640e-33,
            7.51015932e-29]))
```

```
21]: a=pd.Series(a[1])
a.index=x.columns
a
```

```
21]: age          8.580535e-13
sex           4.487078e-24
cp            1.936411e-42
trestbps      6.265162e-04
chol          7.614095e-05
fbs           2.169821e-01
restecg       8.740969e-08
thalach       6.839257e-44
exang         1.812234e-46
oldpeak       1.954551e-48
slope         1.175607e-25
ca            1.884036e-33
thal          7.510159e-29
dtype: float64
```

# Support vector Machine

## Random Forest Regressor

```
In [30]: from sklearn.ensemble import RandomForestClassifier
        from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
```

```
In [31]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.33, random_state=42)
```

```
In [35]: rf = RandomForestClassifier()
        param_grid = {
            'n_estimators': [5, 10, 15],
            'max_depth': [15, 20, 25],
            'min_samples_split': [5, 7, 8]
        }
```

```
In [36]: grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, scoring='accuracy')
```

```
In [37]: grid_search.fit(x_train, y_train)
```

```
Out[37]: GridSearchCV(cv=5, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [15, 20, 25],
                                'min_samples_split': [5, 7, 8],
                                'n_estimators': [5, 10, 15]},
                    scoring='accuracy')
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [38]: grid_search.best_params_
```

```
Out[38]: {'max_depth': 15, 'min_samples_split': 5, 'n_estimators': 15}
```

```
In [40]: best_rf = grid_search.best_estimator_
```

```
In [41]: best_rf
```

```
Out[41]: RandomForestClassifier(max_depth=15, min_samples_split=5, n_estimators=15)
```

```
In [42]: pred=best_rf.predict(x_test)
accuracy_score(pred,y_test)
```

```
Out[42]: 0.9749216300940439
```

```
In [45]: confusion_matrix(pred,y_test)
```

```
Out[45]: array([[157,  5],
               [ 3, 154]], dtype=int64)
```

```
In [46]: print(classification_report(y_test,pred))
```

	precision	recall	f1-score	support
0	0.97	0.98	0.98	160
1	0.98	0.97	0.97	159
accuracy			0.97	319
macro avg	0.98	0.97	0.97	319
weighted avg	0.97	0.97	0.97	319

**Accuracy Score - 0.9749316**